

MPCA CODING ASSIGNMENT

TEAM MEMBERS:

- 1) M NIRANJAN-PES2UG23CS308
- 2) KISHORE H N – PES2UG23CS278
- 3) KEERTHAN P V – PES2UG23CS272

QUESTION NO.3

CPU task scheduling affects cache locality (keeping frequently used data in cache). Implement a task scheduler that optimizes CPU cache utilization. Use a priority queue to schedule tasks while maximizing cache hits. Expected Outcome: A program that simulates task scheduling with cache-aware strategies. Improved CPU efficiency by reducing cache misses. Comparison of FIFO and LRU scheduling policies

SOLUTION FOR THE GIVEN CODE

1) MAIN CODE:

```
import heapq
import random
import matplotlib.pyplot as plt
from collections import deque, OrderedDict

class Task:
    def __init__(self, task_id, access_frequency):
        self.task_id = task_id
```

```
self.access_frequency = access_frequency
```

```
def __lt__(self, other):
```

```
    return self.access_frequency > other.access_frequency # max-heap behavior
```

```
class CacheFIFO:
```

```
    def __init__(self, capacity):
```

```
        self.cache = deque()
```

```
        self.capacity = capacity
```

```
        self.cache_set = set()
```

```
        self.hits = []
```

```
        self.misses = []
```

```
        self.total_hits = 0
```

```
        self.total_misses = 0
```

```
    def access(self, task_id):
```

```
        if task_id in self.cache_set:
```

```
            self.total_hits += 1
```

```
        else:
```

```
            self.total_misses += 1
```

```
            if len(self.cache) >= self.capacity:
```

```
                removed = self.cache.popleft()
```

```
                self.cache_set.remove(removed)
```

```
self.cache.append(task_id)
self.cache_set.add(task_id)
self.hits.append(self.total_hits)
self.misses.append(self.total_misses)
```

```
class CacheLRU:
```

```
def __init__(self, capacity):
    self.cache = OrderedDict()
    self.capacity = capacity
    self.hits = []
    self.misses = []
    self.total_hits = 0
    self.total_misses = 0

def access(self, task_id):
    if task_id in self.cache:
        self.total_hits += 1
        self.cache.move_to_end(task_id)
    else:
        self.total_misses += 1
        if len(self.cache) >= self.capacity:
            self.cache.popitem(last=False)
        self.cache[task_id] = True
    self.hits.append(self.total_hits)
```

```
self.misses.append(self.total_misses)
```

```
def simulate(task_count=20, cache_size=5,  
scheduler_iterations=100):
```

```
    tasks = [Task(task_id=i, access_frequency=random.randint(1, 10))  
for i in range(task_count)]
```

```
    pq = []
```

```
    for task in tasks:
```

```
        heapq.heappush(pq, task)
```

```
    fifo_cache = CacheFIFO(cache_size)
```

```
    lru_cache = CacheLRU(cache_size)
```

```
    print("Starting cache-aware task scheduling simulation...\n")
```

```
    for _ in range(scheduler_iterations):
```

```
        task = heapq.heappop(pq)
```

```
        task_id = task.task_id
```

```
        fifo_cache.access(task_id)
```

```
        lru_cache.access(task_id)
```

```
        task.access_frequency = max(1, task.access_frequency -  
random.randint(0, 2))
```

```
        heapq.heappush(pq, task)
```

```
print(f"Simulation Complete!\n")
print(f"--- FIFO Cache ---")
print(f"Hits: {fifo_cache.total_hits}")
print(f"Misses: {fifo_cache.total_misses}")
print(f"Hit Ratio: {fifo_cache.total_hits / (fifo_cache.total_hits +
fifo_cache.total_misses):.2f}")
fifo_efficiency = fifo_cache.total_hits * 1 + fifo_cache.total_misses
* 10
```

```
print(f"\n--- LRU Cache ---")
print(f"Hits: {lru_cache.total_hits}")
print(f"Misses: {lru_cache.total_misses}")
print(f"Hit Ratio: {lru_cache.total_hits / (lru_cache.total_hits +
lru_cache.total_misses):.2f}")
lru_efficiency = lru_cache.total_hits * 1 + lru_cache.total_misses *
10
```

```
print("\n--- CPU Efficiency Comparison ---")
print(f"Estimated Time Cost with FIFO: {fifo_efficiency} units
(Lower is Better)")
print(f"Estimated Time Cost with LRU : {lru_efficiency} units (Lower
is Better)")
```

```
steps = list(range(1, scheduler_iterations + 1))
plt.figure(figsize=(12, 6))
```

```
plt.plot(steps, fifo_cache.hits, label='FIFO Hits', linestyle='--',
color='blue')

plt.plot(steps, fifo_cache.misses, label='FIFO Misses', linestyle='-',
color='skyblue')

plt.plot(steps, lru_cache.hits, label='LRU Hits', linestyle='--',
color='green')

plt.plot(steps, lru_cache.misses, label='LRU Misses', linestyle='-',
color='lime')

plt.xlabel('Scheduler Iterations')
plt.ylabel('Cache Events')
plt.title('Cache Hit/Miss Comparison: FIFO vs LRU')
plt.legend()
plt.grid(True)
plt.tight_layout()
plt.show()
```

```
# Run the simulation
simulate()
```

2) SCREENSHOT OF THE EXPECTED OUTCOME

Starting cache-aware task scheduling simulation...

Simulation Complete!

--- FIFO Cache ---

Hits: 38

Misses: 62

Hit Ratio: 0.38

--- LRU Cache ---

Hits: 38

Misses: 62

Hit Ratio: 0.38

--- CPU Efficiency Comparison ---

Estimated Time Cost with FIFO: 658 units (Lower is Better)

Estimated Time Cost with LRU : 658 units (Lower is Better)

